

# Docker背后的内核知识

namespace资源隔离



by 李辰洋

# 概念

分类	系统调用参数	相关内核版本
Mount namespace	CLONE_NEWNS	Linux 2.4.19
UTS namespace	CLONE_NEWUTS	Linux 2.6.19
IPC namespace	CLONE_NEWIPC	Linux 2.6.19
PID namespace	CLONE_NEWPID	Linux 2.6.24
Network namespace	CLONE_NEWNET	Linux 2.6.29
User namespace	CLONE_NEWUSER	Linux 3.8

# 概念

- 系统调用

```
int clone(int (*fn)(void *), void *child_stack, int flags, void *arg);
```

<http://man7.org/linux/man-pages/man2/clone.2.html>

- 查看 /proc/<pid>/ns

```
[cyli@R410:~/Docker-Learning/namespace/c$ ps
  PID TTY          TIME CMD
 13215 pts/0        00:00:00 bash
 13305 pts/0        00:00:00 ps
[cyli@R410:~/Docker-Learning/namespace/c$ ls -l /proc/13215/ns
总用量 0
lrwxrwxrwx 1 cyli cyli 0 9月 10 19:07 cgroup -> cgroup:[4026531835]
lrwxrwxrwx 1 cyli cyli 0 9月 10 19:07 ipc -> ipc:[4026531839]
lrwxrwxrwx 1 cyli cyli 0 9月 10 19:07 mnt -> mnt:[4026531840]
lrwxrwxrwx 1 cyli cyli 0 9月 10 19:07 net -> net:[4026531957]
lrwxrwxrwx 1 cyli cyli 0 9月 10 19:07 pid -> pid:[4026531836]
lrwxrwxrwx 1 cyli cyli 0 9月 10 19:07 user -> user:[4026531837]
lrwxrwxrwx 1 cyli cyli 0 9月 10 19:07 uts -> uts:[4026531838]
```

# UTS namespace

- 核心代码

```
#define STACK_SIZE (1024*1024)

static char container_stack[STACK_SIZE];

char* const container_args[] = { "/bin/bash", NULL };

int container_main(void *args) {
    sethostname("tinylcy", 7);
    execv(container_args[0], container_args);
    return -1;
}

int main(void) {
    int container_pid = clone(container_main, container_stack + STACK_SIZE
                             , SIGCHLD | CLONE_NEWUTS, NULL);

    waitpid(container_pid, NULL, 0);
    return 0;
}
```

# UTS namespace

- 宿主机

```
[cyli@R410:~/Docker-Learning/namespace/c$ hostname  
R410  
[cyli@R410:~/Docker-Learning/namespace/c$ echo $$  
13215
```

- 容器

```
[cyli@R410:~/Docker-Learning/namespace/c$ gcc uts_namespace.c -o uts  
[cyli@R410:~/Docker-Learning/namespace/c$ sudo ./uts  
Parent - start a container.  
Container - inside the container.  
[root@tinylcy:~/Docker-Learning/namespace/c# hostname  
tinylcy  
[root@tinylcy:~/Docker-Learning/namespace/c# echo $$  
13718
```

- pstree -pl

```
└─sshd(1267)──sshd(13092)──sshd(13214)──bash(13215)──sudo(13716)──uts(13717)──bash(13718)──pstree(13737)
```

# IPC namespace

- 核心代码

```
int container_pid = clone(container_main, container_stack + STACK_SIZE  
                        , SIGCHLD | CLONE_NEWIPC, NULL);
```

- 创建消息队列

```
[cyli@R410:~/Docker-Learning/namespaces/c$ ipcmk -Q  
消息队列 id: 65536  
[cyli@R410:~/Docker-Learning/namespaces/c$ ipcs -q  
  
----- 消息队列 -----  
键          msqid      拥有者  权限    已用字节数  消息  
0x762aa85c 65536      cyli    644      0           0
```

# IPC namespace

- 执行 UTS

```
[cyli@R410:~/Docker-Learning/namespace/c$ sudo ./uts
Parent - start a container.
Container - inside the container.
[root@tiny1cy:~/Docker-Learning/namespace/c# ipcs -q

----- 消息队列 -----
键          msqid      拥有者  权限    已用字节数  消息
0x762aa85c 65536      cyli    644      0           0
```

- 执行 IPC

```
[cyli@R410:~/Docker-Learning/namespace/c$ sudo ./ipc
Parent - start a container.
Container - inside the container.
[root@R410:~/Docker-Learning/namespace/c# ipcs -q

----- 消息队列 -----
键          msqid      拥有者  权限    已用字节数  消息
```

# PID & Mount namespace

- 核心代码

```
int container_pid = clone(container_main, container_stack + STACK_SIZE  
                        , SIGCHLD | CLONE_NEWPID, NULL);
```

- 执行 PID

```
[cyli@R410:~/Docker-Learning/namespace/c$ gcc pid_namespace.c -o pid  
[cyli@R410:~/Docker-Learning/namespace/c$ sudo ./pid  
Parent - start a container.  
Container - inside the container.  
Current PID: 1  
[root@R410:~/Docker-Learning/namespace/c# echo $$  
1
```



# PID & Mount namespace

- 核心代码

```
int container_pid = clone(container_main, container_stack + STACK_SIZE  
                        , SIGCHLD | CLONE_NEWNS | CLONE_NEWPID, NULL);
```

```
system("mount -t proc proc /proc"); // http://tldp.org/LDP/lfs/5.0/html/chapter06/proc.html
```

- 执行 Mount 前

```
[cyli@R410:~/Docker-Learning/namespace/c$ gcc mount_namespace.c -o mount  
[cyli@R410:~/Docker-Learning/namespace/c$ ls /proc
```

1	109	1228	1270	131	141	149	1535	1848	20	248	2848	291
10	1099	1236	1271	1314	142	15	156	1856	2126	249	285	292
100	11	1238	128	132	143	150	16	1859	22	25	2852	293
101	1117	1243	1286	1334	144	151	169	1860	23	27	2854	294
102	1173	1245	129	134	1442	1513	170	1861	233	270	2857	295
103	1198	1247	13	135	145	152	171	19	24	271	2859	296
105	12	1249	130	1353	146	1531	18	1929	244	276	286	297
106	1216	125	1305	136	147	1532	1818	1938	245	278	289	2971
107	1218	1251	1306	137	1470	1533	1846	194	246	28	29	2972
108	1223	126	1309	14	148	1534	1847	2	247	283	290	3

# PID & Mount namespace

- 核心代码

```
int container_pid = clone(container_main, container_stack + STACK_SIZE  
                        , SIGCHLD | CLONE_NEWNS | CLONE_NEWPID, NULL);
```

```
system("mount -t proc proc /proc"); // http://tldp.org/LDP/lfs/5.0/html/chapter06/proc.html
```

- 执行 Mount 后

```
[cyli@R410:~/Docker-Learning/namespace/c$ sudo ./mount  
Parent - start a container.  
Container - inside the container.  
[root@R410:~/Docker-Learning/namespace/c# ls /proc  
1          bus          cpuinfo    dma         filesystems ioports  
13         cgroups  crypto     driver      fs           ipmi  
acpi       cmdline  devices    execdomains interrupts   irq  
buddyinfo consoles diskstats fb          iomem       kallsyms
```

# User namespace

- 设置CLONE\_NEWUSER

```
[cyli@R410:~/Docker-Learning/namespace/c$ gcc user_namespace.c -o user
[cyli@R410:~/Docker-Learning/namespace/c$ sudo ./user
Parent: eUID = 0, eGID = 0, UID = 0, GID = 0
Parent: start a container.
Parent: user/group mapping done.
Container PID = [5979] - inside the container
Container eUID = 65534, eGID = 65534, UID = 65534, GID = 65534
[nobody@R410:~/Docker-Learning/namespace/c$ id
uid=65534(nobody) gid=65534(nogroup) 组=65534(nogroup)
```

- UID / GID 映射
  - /proc/<pid>/uid\_map
  - /proc/<pid>/gid\_map
  - ID-inside-ns ID-outside-ns length

# User namespace

- 核心代码

```
int pipefd[2];

int main(void) {
    const int gid = getgid(), uid = getuid();

    pipe(pipefd);

    int container_pid = clone(container_main, container_stack + STACK_SIZE
                              , CLONE_NEWUSER | SIGCHLD, NULL);

    set_uid_map(container_pid, 0, uid, 1);
    set_gid_map(container_pid, 0, gid, 1);

    close(pipefd[1]);

    waitpid(container_pid, NULL, 0);

    return 0;
}
```

# User namespace

- 核心代码

```
int container_main(void *args) {  
    char ch;  
    close(pipefd[1]);  
    read(pipefd[0], &ch, 1);  
    execv(container_args[0], container_args);  
    return -1;  
}
```

- 执行User

```
[cyli@R410:~/Docker-Learning/namespace/c$ gcc user_namespace.c -o user  
[cyli@R410:~/Docker-Learning/namespace/c$ ./user  
Parent: eUID = 1002, eGID = 1002, UID = 1002, GID = 1002  
Parent: start a container.  
Parent: user/group mapping done.  
Container PID = [6203] - inside the container  
Container eUID = 0, eGID = 65534, UID = 0, GID = 65534  
[root@R410:~/Docker-Learning/namespace/c# id  
uid=0(root) gid=65534(nogroup) 组=65534(nogroup)
```

# Network namespace

- 创建 namespace

```
[cyli@R410:~$ sudo ip netns add tinyncy_ns  
[cyli@R410:~$ sudo ip netns exec tinyncy_ns ip link list  
1: lo: <LOOPBACK> mtu 65536 qdisc noop state DOWN mode DEFAULT group default qlen 1  
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
```

- 启动 loopback 设备

```
[cyli@R410:~$ sudo ip netns exec tinyncy_ns ping 127.0.0.1  
connect: Network is unreachable  
[cyli@R410:~$ sudo ip netns exec tinyncy_ns ip link set dev lo up  
[cyli@R410:~$ sudo ip netns exec tinyncy_ns ping 127.0.0.1  
PING 127.0.0.1 (127.0.0.1) 56(84) bytes of data.  
64 bytes from 127.0.0.1: icmp_seq=1 ttl=64 time=0.070 ms  
64 bytes from 127.0.0.1: icmp_seq=2 ttl=64 time=0.028 ms
```

# Network namespace

- 创建 VETH(Virtual ETHernet) Pair

```
[cyli@R410:~$ sudo ip link add veth0 type veth peer name veth1  
[cyli@R410:~$ sudo ip link set veth1 netns tinyncy_ns  
[cyli@R410:~$ sudo ip netns exec tinyncy_ns ifconfig veth1 10.1.1.1/24 up  
[cyli@R410:~$ sudo ifconfig veth0 10.1.1.2/24 up
```

- 测试 (VETH)

```
[cyli@R410:~$ ip link list | grep veth0  
19: veth0@if18: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP mode DEFAULT group default qlen 1000  
[cyli@R410:~$ sudo ip netns exec tinyncy_ns ip link list  
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1  
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00  
18: veth1@if19: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP mode DEFAULT group default qlen 1000  
    link/ether f6:db:ea:b8:7a:39 brd ff:ff:ff:ff:ff:ff link-netnsid 0
```

# Network namespace

- 测试 (ping)

```
[cyli@R410:~$ ping 10.1.1.1
PING 10.1.1.1 (10.1.1.1) 56(84) bytes of data.
64 bytes from 10.1.1.1: icmp_seq=1 ttl=64 time=0.054 ms
64 bytes from 10.1.1.1: icmp_seq=2 ttl=64 time=0.027 ms
```

```
[cyli@R410:~$ sudo ip netns exec tinyncy_ns ping 10.1.1.2
PING 10.1.1.2 (10.1.1.2) 56(84) bytes of data.
64 bytes from 10.1.1.2: icmp_seq=1 ttl=64 time=0.080 ms
64 bytes from 10.1.1.2: icmp_seq=2 ttl=64 time=0.033 ms
```

- 测试 (路由表)

```
[cyli@R410:~$ route
内核 IP 路由表
目标          网关          子网掩码      标志  跃点  引用  使用  接口
default       222.201.145.254 0.0.0.0      UG    0     0     0    eno1
10.1.1.0       *             255.255.255.0 U     0     0     0    veth0
172.17.0.0     *             255.255.0.0  U     0     0     0    docker0
192.168.10.0   *             255.255.255.0 U     0     0     0    lxcbr0
192.168.122.0  *             255.255.255.0 U     0     0     0    virbr0
222.201.145.128 *            255.255.255.128 U     0     0     0    eno1
222.201.145.128 *            255.255.255.128 U     0     0     0    eno2
[cyli@R410:~$ sudo ip netns exec tinyncy_ns route
内核 IP 路由表
目标          网关          子网掩码      标志  跃点  引用  使用  接口
10.1.1.0       *             255.255.255.0 U     0     0     0    veth1
```



THANKS

<https://github.com/scut-ccmdp-lab/Docker-Learning/tree/master/namespace>